8B NIM HTML-Only Judge Matrix

Side-by-side comparison of Nemotron Ultra, Claude Sonnet 4.6, and Kimi K2.6 as judges for the HTML-only NIM grounded dataset experiment.

Generated 2026-06-11
Dataset: `nim_curated_html_only`
Completion run: `20260610t032644z`
Kimi eval run: `20260610t204500z`

Scope
8B base plus HTML-only e5 LoRA adapters `r16` and `r32`.

Judges
`nvidia/nvidia/nemotron-3-ultra`, `azure/anthropic/claude-sonnet-4-6`, and `nvidia/moonshotai/kimi-k2.6`.

Rows
482 context-baked eval rows. Kimi completed all 4,338 judge rows with zero permanent failures.

49B comparator
Matching HTML-only 49B completions were generated; older non-HTML 49B rows were not reused.

Phase 1: Single-Axis Scores

Score means on 1-5 axes. Composite is the unweighted mean of accuracy, completeness, faithfulness, and clarity.

Higher is better | All rows complete

Target	Judge	Accuracy	Completeness	Faithfulness	Clarity	Composite	Composite Delta vs Base
8B base	Nemotron Ultra	4.133	3.981	4.361	4.876	4.338	0.000
8B HTML-only r16	Nemotron Ultra	4.162	3.834	4.413	4.907	4.329	**-0.009**
8B HTML-only r32	Nemotron Ultra	4.178	3.840	4.394	4.911	4.331	**-0.007**
8B base	Claude Sonnet 4.6	4.127	3.369	4.400	4.434	4.082	0.000
8B HTML-only r16	Claude Sonnet 4.6	4.210	3.220	4.446	4.506	4.095	**+0.012**
8B HTML-only r32	Claude Sonnet 4.6	4.232	3.100	4.463	4.546	4.085	**+0.003**
8B base	Kimi K2.6	4.139	3.571	4.338	4.668	4.179	0.000
8B HTML-only r16	Kimi K2.6	4.176	3.554	4.390	4.718	4.210	**+0.031**
8B HTML-only r32	Kimi K2.6	4.197	3.454	4.444	4.797	4.223	**+0.044**

Phase 1: Composite Matrix

Compact composite comparison by target and judge.

Target	Nemotron Ultra	Claude Sonnet 4.6	Kimi K2.6	Composite Rank by Judge
8B base	4.338	4.082	4.179	Ultra: 1st, Sonnet: 3rd, Kimi: 3rd
8B HTML-only r16	4.329	4.095	4.210	Ultra: 3rd, Sonnet: 1st, Kimi: 2nd
8B HTML-only r32	4.331	4.085	4.223	Ultra: 2nd, Sonnet: 2nd, Kimi: 1st

Phase 2: 8B Pairwise Matrix

Within-family pairwise comparisons. Win columns follow the pair label order. Right non-loss means right wins plus ties.

Non-loss = wins + ties

Kimi: zero failures

Pair	Judge	Rows	Failed	Left Wins	Right Wins	Ties	Right Non-Loss	Raw Win-Lean
8B base vs r16	Nemotron Ultra	479	3	164	180	135	65.8%	**r16**
8B base vs r16	Claude Sonnet 4.6	482	0	174	145	163	63.9%	**base**
8B base vs r16	Kimi K2.6	482	0	181	214	87	62.4%	**r16**
8B base vs r32	Nemotron Ultra	482	0	156	173	153	67.6%	**r32**
8B base vs r32	Claude Sonnet 4.6	482	0	192	132	158	60.2%	**base**
8B base vs r32	Kimi K2.6	482	0	203	184	95	57.9%	**base**
r16 vs r32	Nemotron Ultra	482	0	149	139	194	69.1%	**r16**
r16 vs r32	Claude Sonnet 4.6	482	0	156	81	245	67.6%	**r16**
r16 vs r32	Kimi K2.6	482	0	186	132	164	61.4%	**r16**

Phase 3: 49B Comparator Matrix

49B is the left side in every pair. The key 8B metric is non-loss: 8B wins plus ties.

Higher 8B non-loss is better All rows complete

Pair	Judge	49B Wins	8B Wins	Ties	8B Non-Loss	49B Non-Loss	Rows
49B vs 8B base	Nemotron Ultra	325	64	93	32.6%	86.7%	482
49B vs 8B base	Claude Sonnet 4.6	303	80	99	37.1%	83.4%	482
49B vs 8B base	Kimi K2.6	338	82	62	29.9%	83.0%	482
49B vs 8B HTML-only r16	Nemotron Ultra	314	98	70	34.9%	79.7%	482
49B vs 8B HTML-only r16	Claude Sonnet 4.6	309	104	69	35.9%	78.4%	482
49B vs 8B HTML-only r16	Kimi K2.6	308	103	71	36.1%	78.6%	482
49B vs 8B HTML-only r32	Nemotron Ultra	316	87	79	34.4%	82.0%	482
49B vs 8B HTML-only r32	Claude Sonnet 4.6	330	109	43	31.5%	77.4%	482
49B vs 8B HTML-only r32	Kimi K2.6	323	91	68	33.0%	81.1%	482

Kimi Token Rollup

Token accounting from the completed Kimi run. Combined tokens include target completion context plus judge scoring tokens.

Phase	Scope	Rows	Judge Tokens	Combined Tokens
Phase 1	8B base single-axis	482	2,574,966	3,801,135
Phase 1	8B r16 single-axis	482	2,492,448	3,681,869
Phase 1	8B r32 single-axis	482	2,520,267	3,713,305
Phase 2	8B pairwise matrix	1,446	14,769,611	21,986,867
Phase 3	49B vs 8B pairwise matrix	1,446	15,508,588	23,957,467
Total	**All Kimi eval phases**	**4,338**	**37,865,880**	**57,140,643**

Interpretation Matrix

Compact readout of what changes after adding Kimi as an independent judge.

Question	Ultra Read	Sonnet Read	Kimi Read	Practical Takeaway
Did LoRA improve single-axis faithfulness?	Yes. r16 improves from 4.361 to 4.413; r32 to 4.394.	Yes. r16 improves from 4.400 to 4.446; r32 to 4.463.	Yes. r16 improves from 4.338 to 4.390; r32 to 4.444.	All three judges agree the adapters improve faithfulness.
Did LoRA improve pairwise preference over base?	Yes for both r16 and r32.	No. Sonnet prefers base over both adapters on raw wins.	Mixed. r16 beats base; base beats r32.	The strongest cross-judge signal is r16, not r32.
Which rank is stronger?	r16 narrowly beats r32 pairwise.	r16 clearly beats r32 pairwise.	r16 clearly beats r32 pairwise.	All three judges prefer r16 over r32 in direct pairwise comparison.
Did 8B approach 49B on Phase 3?	Not parity. Best 8B non-loss is r16 at 34.9%.	Not parity. Best 8B non-loss is base at 37.1%, r16 at 35.9%.	Not parity. Best 8B non-loss is r16 at 36.1%.	49B remains clearly ahead, but Kimi reinforces r16 as the best LoRA candidate for follow-up.
Does adding Kimi change the recommendation?	Kimi is closer to Ultra than Sonnet on pairwise LoRA preference, but less favorable to r32. It supports r16 as the most defensible adapter.			Use r16 as the default 8B HTML-only LoRA candidate; treat r32 as an axis-score improvement, not the pairwise winner.