8B NIM HTML-Only Judge Matrix

Side-by-side comparison of Nemotron Ultra and Claude Sonnet 4.6 as judges for the HTML-only NIM grounded dataset experiment.

Generated 2026-06-10
Dataset: `nim_curated_html_only`
Completion run: `20260610t032644z`

Scope
8B base plus HTML-only e5 LoRA adapters `r16` and `r32`.

Judges
`nvidia/nvidia/nemotron-3-ultra` and `azure/anthropic/claude-sonnet-4-6`.

Rows
482 context-baked eval rows. Claude has zero row failures across all phases.

49B comparator
Matching HTML-only 49B completions were generated; older non-HTML 49B rows were not reused.

Phase 1: Single-Axis Matrix

Score means on 1-5 axes. Delta columns are Sonnet minus Ultra.

Higher is better

Composite = mean of four axes

Target	Nemotron Ultra					Clauc	
	Accuracy	Completeness	Faithfulness	Clarity	Composite	Accuracy	Completeness
8B base	4.133	3.981	4.361	4.876	4.338	4.127	3.369
8B HTML-only r16	4.162	3.834	4.413	4.907	4.329	4.210	3.220
8B HTML-only r32	4.178	3.840	4.394	4.911	4.331	4.232	3.100

Phase 2: 8B Pairwise Matrix

Within-family pairwise comparisons. Win columns follow the pair label order.

Non-loss = wins + ties

Ultra base-vs-r16 has 3 failed rows

Pair	Rows	Failed	Nemotron Ultra						
			Left Wins	Right Wins	Ties	Right Non-Loss	Win-Lean	Left Wins	R
8B base vs r16	479 / 482	3 / 0	164	180	135	65.8%	**r16**	174	
8B base vs r32	482	0	156	173	153	67.6%	**r32**	192	
r16 vs r32	482	0	149	139	194	69.1%	**r16**	156	

Phase 3: 49B Comparator Matrix

49B is the left side in every pair. The key metric is 8B non-loss: 8B wins plus ties.

Higher 8B non-loss is better for adaptation

All rows complete, zero failures

Pair	Nemotron Ultra						Cl		
	49B Wins	8B Wins	Ties	8B Non-Loss	49B Non-Loss	Rows	49B Wins	8B Wins	Ties
49B vs 8B base	325	64	93	32.6%	86.7%	482	303	80	99
49B vs 8B HTML-only r16	314	98	70	34.9%	79.7%	482	309	104	69
49B vs 8B HTML-only r32	316	87	79	34.4%	82.0%	482	330	109	43

Interpretation Matrix

Compact readout of what the two judges agree and disagree on.

Question	Ultra Read	Sonnet Read	Pra
Did LoRA improve single-axis faithfulness?	Yes. r16 improved from 4.361 to 4.413; r32 to 4.394.	Yes. r16 improved from 4.400 to 4.446; r32 to 4.463.	Bo fai
Did LoRA improve pairwise preference over base?	Yes. r16 and r32 both beat base on raw wins.	No. Sonnet prefers base over both adapters on raw wins.	Juc axi
Which rank is stronger?	r16 narrowly beats r32 by raw wins; r32 slightly leads single-axis accuracy/completeness.	r16 clearly beats r32 pairwise; r32 leads single-axis accuracy/faithfulness.	r16 iso
Did 8B approach 49B on Phase 3?	Not parity. Best 8B non-loss is r16 at 34.9%.	Not parity. Best 8B non-loss is base at 37.1%, r16 at 35.9%.	HT ga jus